HIVE:
Question 1.

1. FInd the airports that are only listed as sources but not as destinations in the routes table.

```
SELECT
      DISTINCT(src.name)
FROM airport src
LEFT JOIN routes r
ON src.airport_id = r.src_airport_id
WHERE src.airport_id IS NOT NULL;
```

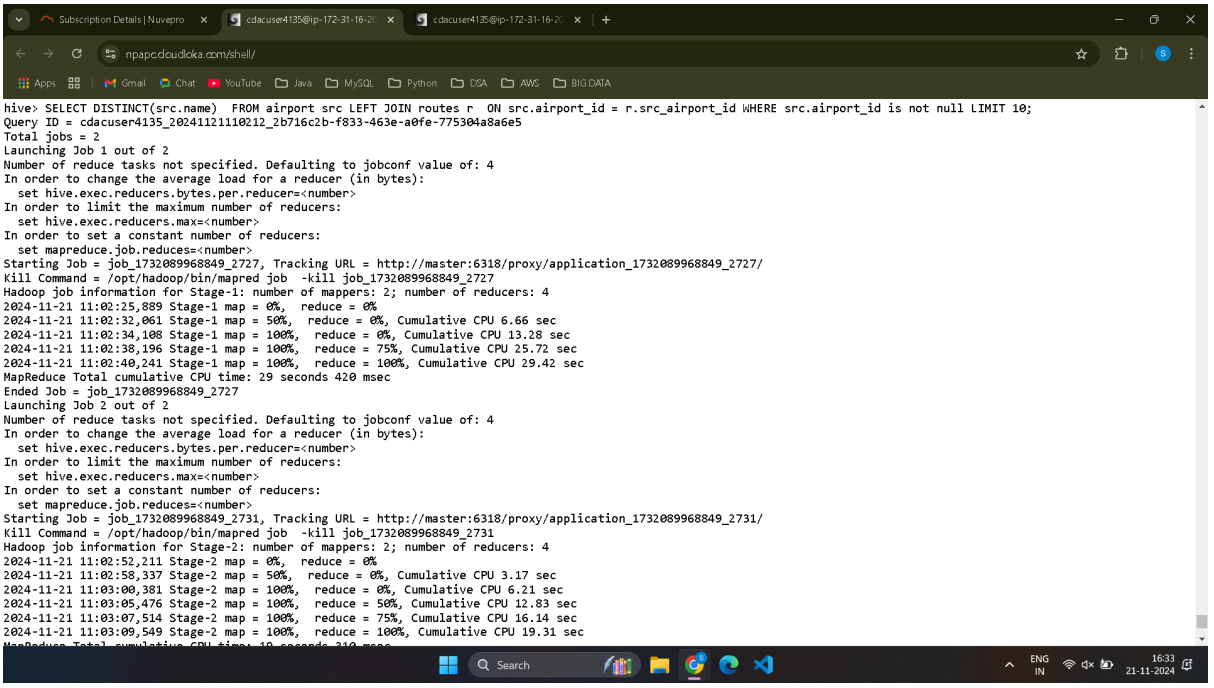

```
hive> SELECT DISTINCT(src.name)  FROM airport src LEFT JOIN routes r  ON src.airport_id = r.src_airport_id WHERE src.airport_id is not null LIMIT 10;
Query ID = cdacuser4135_20241121110212_2b716c2b-f833-463e-a0fe-775304a8a6e5
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Defaulting to jobconf value of: 4
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1732089968849_2727, Tracking URL = http://master:6318/proxy/application_1732089968849_2727/
Kill Command = /opt/hadoop/bin/mapred job  -kill job_1732089968849_2727
Hadoop job information for Stage-1: number of mappers: 2; number of reducers: 4
2024-11-21 11:02:25,889 Stage-1 map = 0%,  reduce = 0%
2024-11-21 11:02:32,061 Stage-1 map = 50%,  reduce = 0%, Cumulative CPU 6.66 sec
2024-11-21 11:02:34,108 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 13.28 sec
2024-11-21 11:02:38,196 Stage-1 map = 100%,  reduce = 75%, Cumulative CPU 25.72 sec
2024-11-21 11:02:40,241 Stage-1 map = 100%,  reduce = 100%, Cumulative CPU 29.42 sec
MapReduce Total cumulative CPU time: 29 seconds 420 msec
Ended Job = job_1732089968849_2727
Launching Job 2 out of 2
Number of reduce tasks not specified. Defaulting to jobconf value of: 4
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1732089968849_2731, Tracking URL = http://master:6318/proxy/application_1732089968849_2731/
Kill Command = /opt/hadoop/bin/mapred job  -kill job_1732089968849_2731
Hadoop job information for Stage-2: number of mappers: 2; number of reducers: 4
2024-11-21 11:02:52,211 Stage-2 map = 0%,  reduce = 0%
2024-11-21 11:02:58,337 Stage-2 map = 50%,  reduce = 0%, Cumulative CPU 3.17 sec
2024-11-21 11:03:00,381 Stage-2 map = 100%,  reduce = 0%, Cumulative CPU 6.21 sec
2024-11-21 11:03:05,476 Stage-2 map = 100%,  reduce = 50%, Cumulative CPU 12.83 sec
2024-11-21 11:03:07,514 Stage-2 map = 100%,  reduce = 75%, Cumulative CPU 16.14 sec
2024-11-21 11:03:09,549 Stage-2 map = 100%,  reduce = 100%, Cumulative CPU 19.31 sec
```

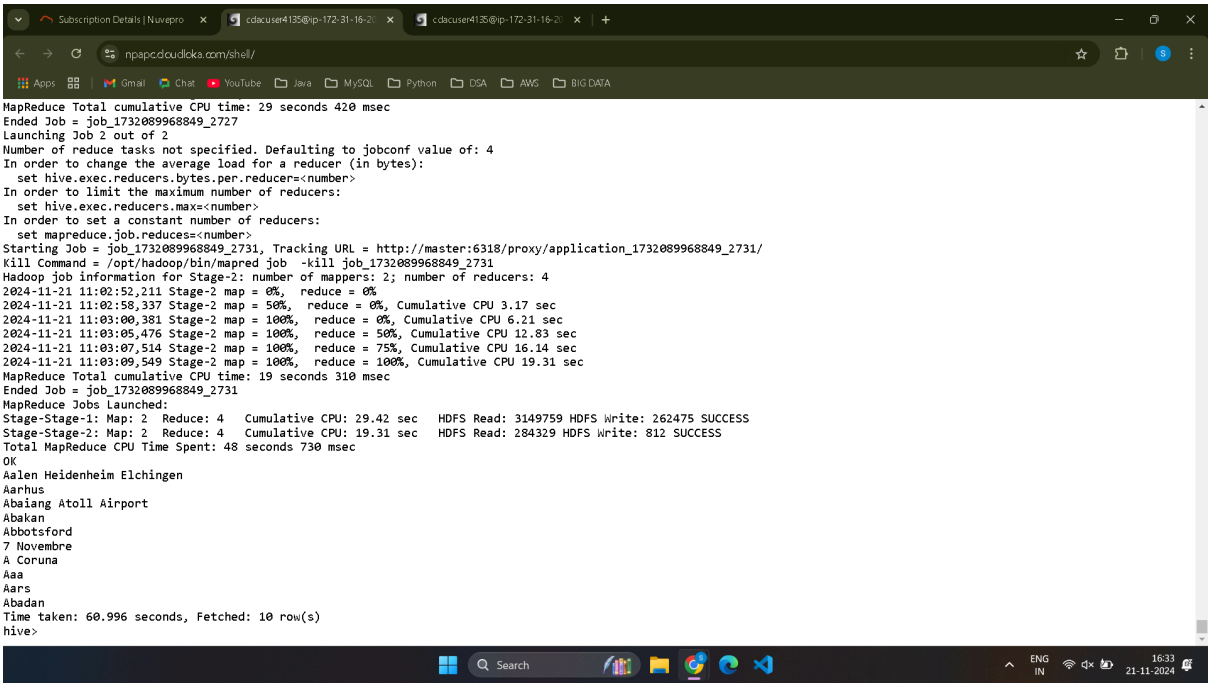

```
MapReduce Total cumulative CPU time: 29 seconds 420 msec
Ended Job = job_1732089968849_2727
Launching Job 2 out of 2
Number of reduce tasks not specified. Defaulting to jobconf value of: 4
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1732089968849_2731, Tracking URL = http://master:6318/proxy/application_1732089968849_2731/
Kill Command = /opt/hadoop/bin/mapred job  -kill job_1732089968849_2731
Hadoop job information for Stage-2: number of mappers: 2; number of reducers: 4
2024-11-21 11:02:52,211 Stage-2 map = 0%,  reduce = 0%
2024-11-21 11:02:58,337 Stage-2 map = 50%,  reduce = 0%, Cumulative CPU 3.17 sec
2024-11-21 11:03:00,381 Stage-2 map = 100%,  reduce = 0%, Cumulative CPU 6.21 sec
2024-11-21 11:03:05,476 Stage-2 map = 100%,  reduce = 50%, Cumulative CPU 12.83 sec
2024-11-21 11:03:07,514 Stage-2 map = 100%,  reduce = 75%, Cumulative CPU 16.14 sec
2024-11-21 11:03:09,549 Stage-2 map = 100%,  reduce = 100%, Cumulative CPU 19.31 sec
MapReduce Total cumulative CPU time: 19 seconds 310 msec
Ended Job = job_1732089968849_2731
MapReduce Jobs Launched:
Stage-Stage-1: Map: 2  Reduce: 4   Cumulative CPU: 29.42 sec   HDFS Read: 3149759 HDFS Write: 262475 SUCCESS
Stage-Stage-2: Map: 2  Reduce: 4   Cumulative CPU: 19.31 sec   HDFS Read: 284329 HDFS Write: 812 SUCCESS
Total MapReduce CPU Time Spent: 48 seconds 730 msec
OK
Aalen Heidenheim Elchingen
Aarhus
Abaiang Atoll Airport
Abakan
Abbotsford
7 Novembre
A Coruna
Aaa
Aars
Abadan
Time taken: 60.996 seconds, Fetched: 10 row(s)
hive>
```

2. Determine the top 3 airlines that operate on the highest number of distinct routes.

```
SELECT
      airline_id, count(*)
FROM routes
GROUP BY airline_id
ORDER BY count(*) DESC
LIMIT 3;
```

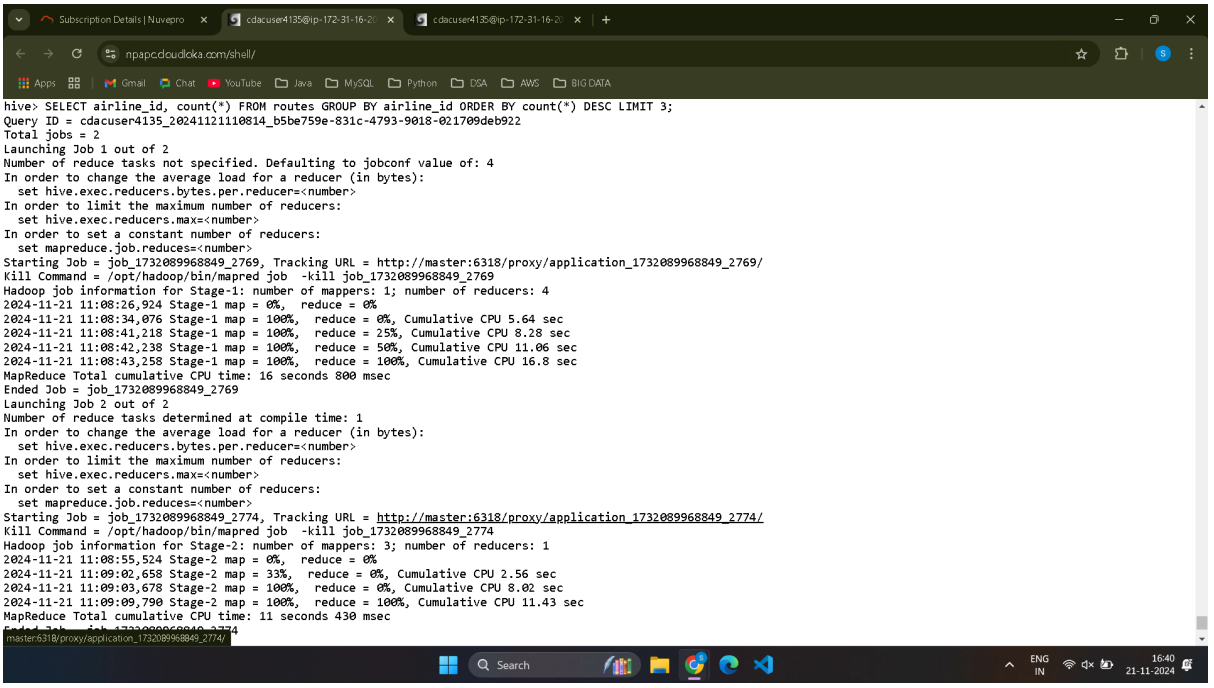

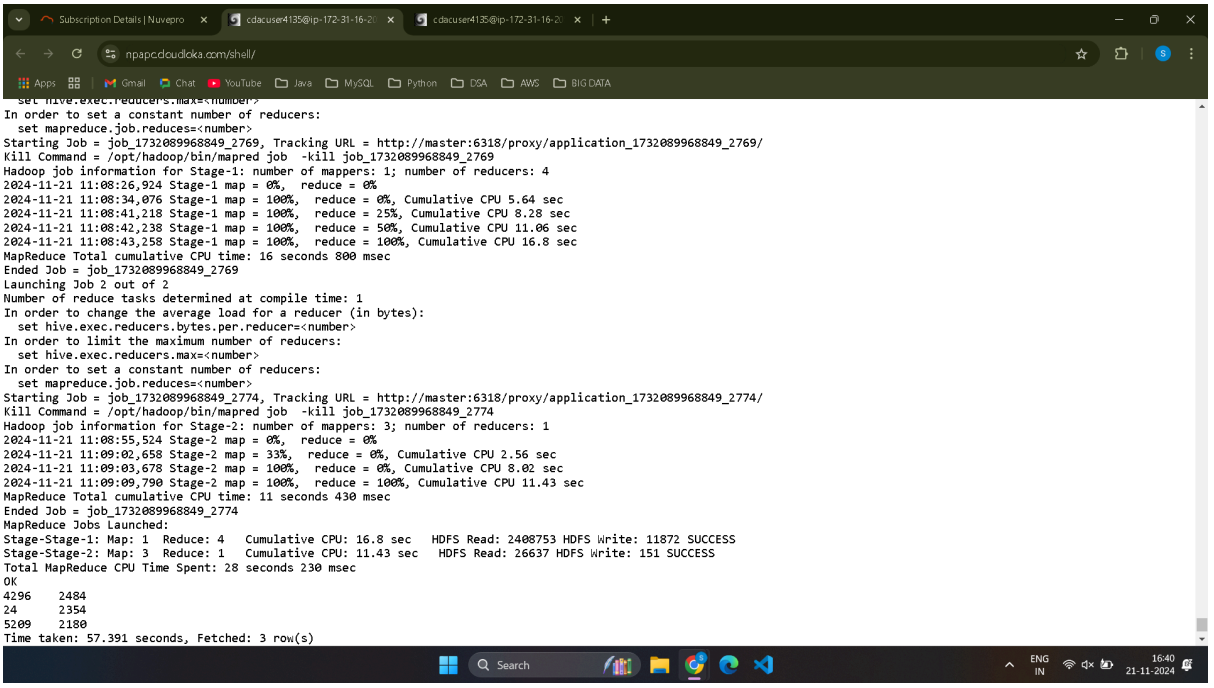

3. Find the total number of distinct aircraft types (Equipment) used in the routes table.

```
SELECT
      COUNT(*)
FROM (SELECT
      COUNT(*)
FROM routes
GROUP BY equipment)abc;
```

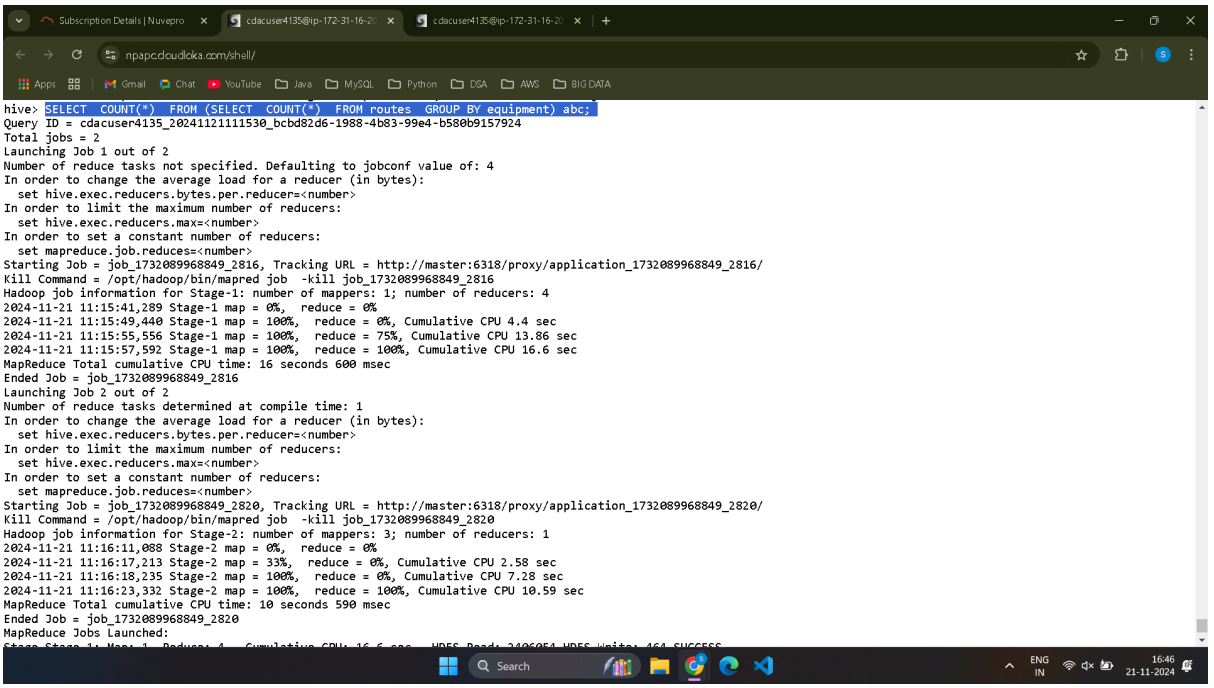

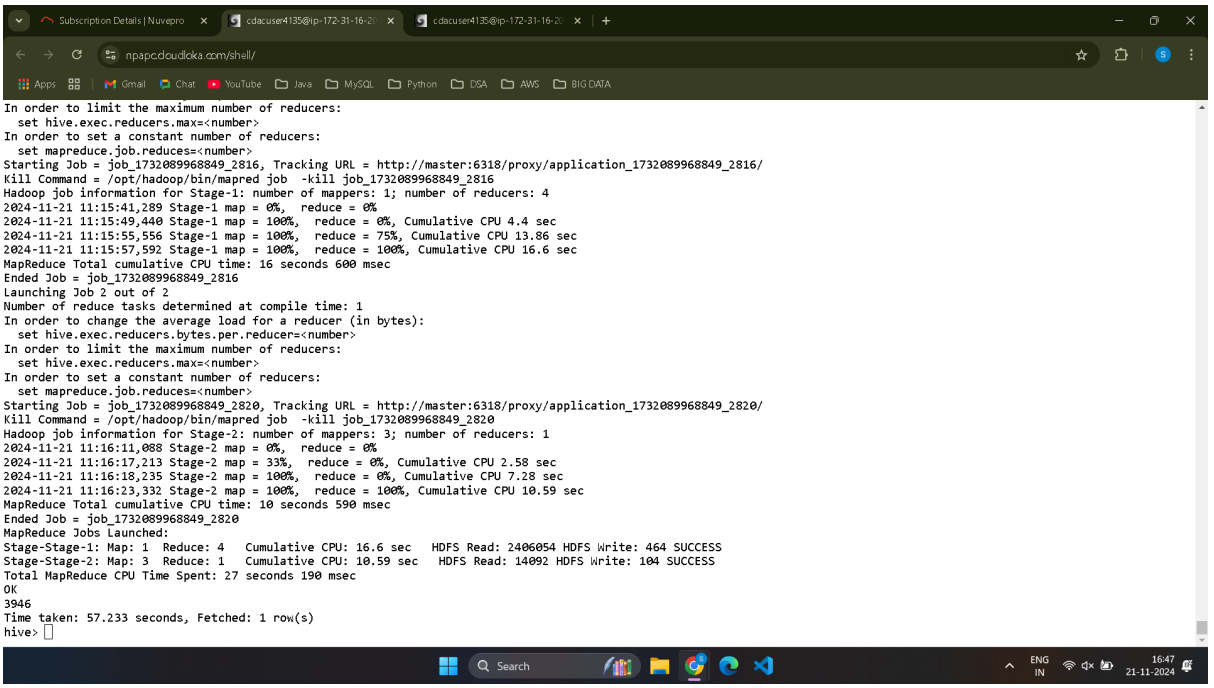

Question 2.

1. Write an SQL query to create a partitioned table named routes_partitioned to store route details, partitioned by the Destination_Airport Column

```
CREATE TABLE routes_partioned
      (STRING airline_iata,
      INT airline_id,
      STRING src_airport_iata,
      INT  src_airport_id,
      INT dest_airport_id,
      STRING codeshare,
      INT stops,
      STRING equipment)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ‘,’
STORED AS TEXTFILE
PARTITIONED BY(STRING dest_airport_iata)
‘/user/cdacuser4135/routes.csv’
```

2. Insert data into the routes_partitioned table for a specific Destination_Airport
3. To query data for a specific Destination_Airport
4. Write a query to display all available partitions in the routes_partitioned table.

SPARK:

```
>>> Airlines = sc.textFile("/user/cdacuser4135/airlines.csv")
>>> Airlines.take(5)
>>> header = Airlines.first()
>>> Airlines = Airlines.filter(lambda line: line != header)
>>> Airlines.take(5)
```

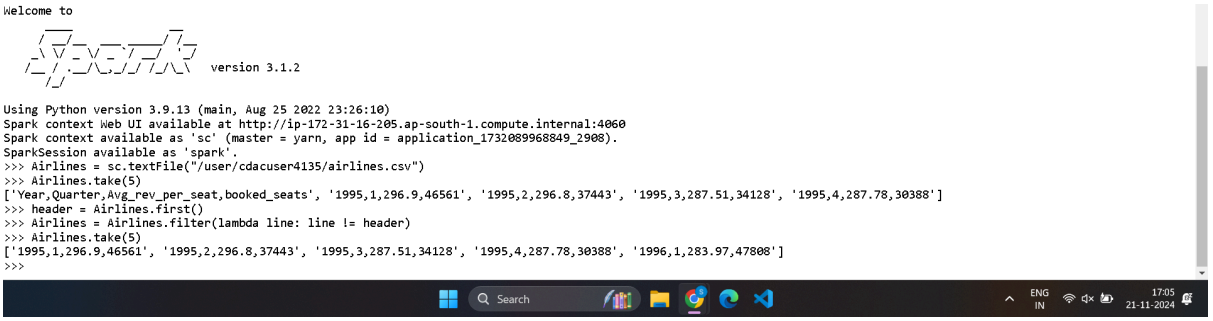

Question 1.

1. Count the number of rows where the total no. of booked seats is between 20,000 and 50,000

2. Find and list all the distinct quarters and years (combined as tuples) present in the dataset. Ensure no duplicates are included in the result.

```
>>> Distinct_Quarters_Years = Airlines_split.map(lambda a: (a[0], a[1]))
>>> Distinct_Quarters_Years.take(20)
```

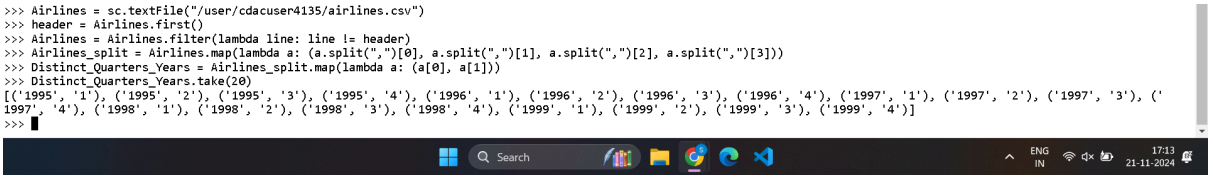

Question 2.

```
>>> AirlinesDF = spark.read.csv("/user/cdacuser4135/airlines.csv", header = True, inferSchema = True)
>>> AirlinesDF.show()
>>> from pyspark.sql import functions as F
```

```
>>> AirlinesDF = spark.read.csv("/user/cdacuser4135/airlines.csv", header = True, inferSchema = True)
>>> AirlinesDF.show()
+----+-------+---------------+------------+
|Year|Quarter|Avg_rev_per_seat|booked_seats|
+----+-------+---------------+------------+
|1995|      1|          296.9|       46561|
|1995|      2|          296.8|       37443|
|1995|      3|         287.51|       34128|
|1995|      4|         287.78|       30388|
|1996|      1|         283.97|       47808|
|1996|      2|         275.78|       43020|
|1996|      3|         269.49|       38952|
|1996|      4|         278.33|       37443|
|1997|      1|          283.4|       35067|
|1997|      2|         289.44|       46565|
|1997|      3|         282.27|       38886|
|1997|      4|         293.51|       37454|
|1998|      1|         304.74|       31315|
|1998|      2|         300.97|       30852|
|1998|      3|         315.25|       38118|
|1998|      4|         316.18|       35393|
|1999|      1|         331.74|       47453|
|1999|      2|         329.34|       38243|
|1999|      3|         317.22|       33048|
|1999|      4|         317.93|       31256|
+----+-------+---------------+------------+
only showing top 20 rows

>>> from pyspark.sql import functions as F
```

1. Find the minimum, maximum and average of the total number of booked seats.

```
>>> Max_Booked_Seats = AirlinesDF.agg(F.max("booked_seats"))
>>> Max_Booked_Seats.show()
+------------------+
|max(booked_seats)|
+------------------+
|             49678|
+------------------+

>>> Min_Booked_Seats = AirlinesDF.agg(F.min("booked_seats"))
>>> Min_Booked_Seats.show()
+------------------+
|min(booked_seats)|
+------------------+
|             30103|
+------------------+

>>> Avg_Booked_Seats = AirlinesDF.agg(F.avg("booked_seats"))
>>> Avg_Booked_Seats.show()
+------------------+
|avg(booked_seats)|
+------------------+
|39640.70238095238|
+------------------+
```

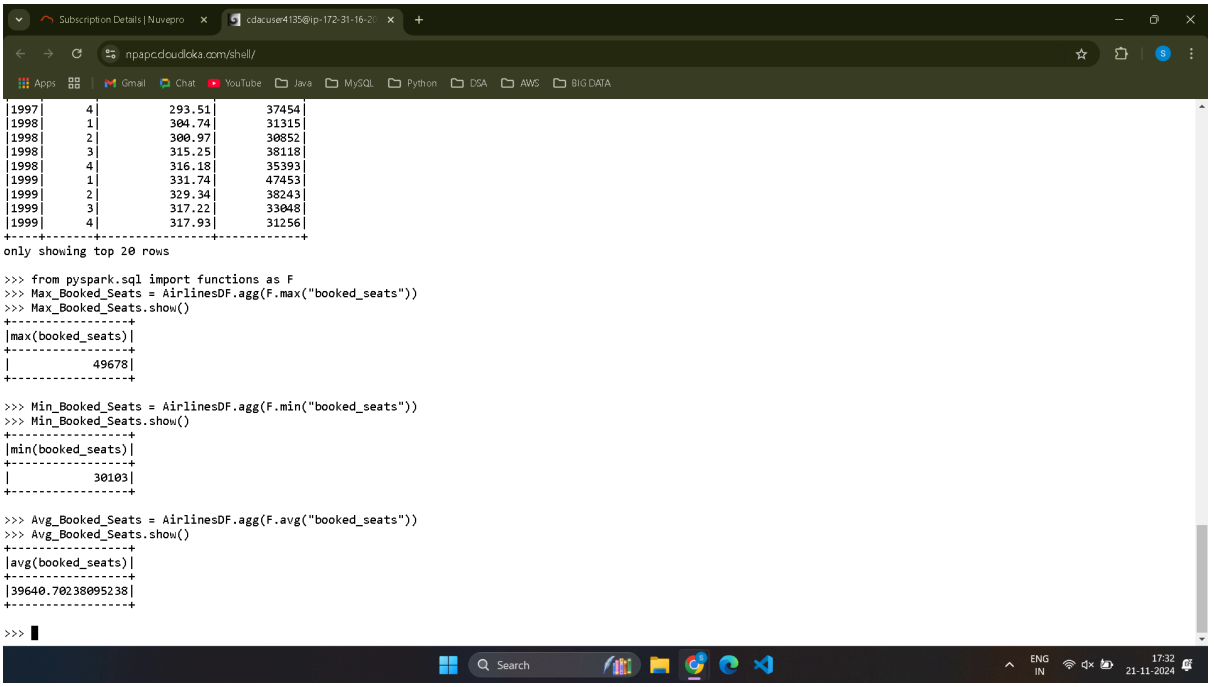

2. Count the number of rows where the average revenue per seat is less than $290.

3. Find the average number of booked seats grouped by quarter.

```
>>> Avg_booked_seat_by_quarter =
AirlinesDF.groupBy("quarter").agg(F.mean("booked_seats"))
>>> Avg_booked_seat_by_quarter.show()
```

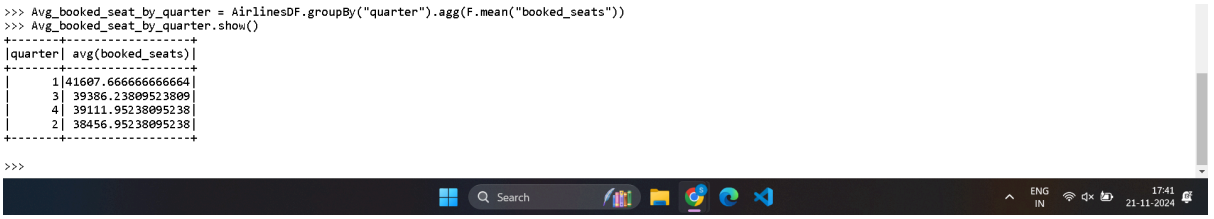

4. List all distinct years in the dataset along with the count of rows for each year.

```
>>> Distinct_yrs_count =
AirlinesDF.groupBy("Year").agg(F.count("Year"))
>>> Distinct_yrs_count.show()
```

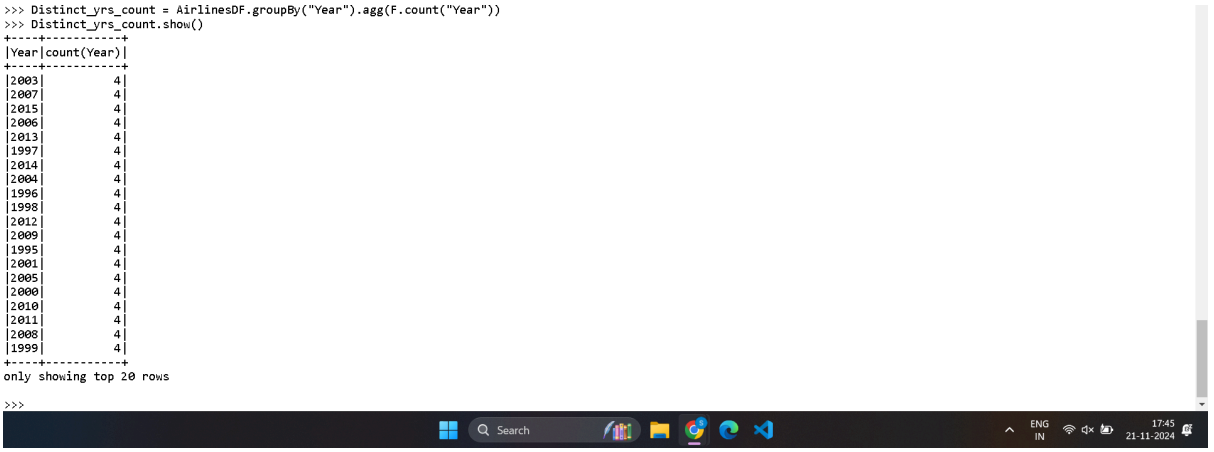

5. Find the quarter with the highest total revenue across the years and its value.

```
>>> Airlines = sc.textFile('/user/cdacuser4135/airlines.csv')
>>> header = Airlines.first()
>>> Airlines = Airlines.filter(lambda line: line != header)
>>> Airlines_split = Airlines.map(lambda a: (a.split(",")[1],
float(a.split(",")[2]), int(a.split(",")[3])))
>>> Total_rev = Airlines.map(lambda a: (a[0],
float(a[1])*int(a[2])))
>>> Tot_rev_by_quarter = Total_rev.reduceByKey(lambda a,b :
a+b).sortBy(lambda a: -a[1])
>>> Tot_rev_by_quarter.take(1)
[('1', 1620.0)]
```